SQL Bits - the great data heist Manchester 2019

An R primer for SQL folks

Thomas Hütter

An R primer for SQL folks

Thomas Hütter, Diplom-Betriebswirt

- Application developer, consultant, accidental DBA, author
- Worked at consultancies, ISVs, end user companies
- Speaker at SQL events around Europe
- SQL Server > 6.5, Dynamics Nav > 3.01, R > 3.1.2

 @DerFredo https://twitter.com/DerFredo

 de.linkedin.com/in/derfredo

 www.xing.com/profile/Thomas_Huetter

Agenda

- History: what is R, how did R come to be,
  what does the R ecosystem look like today

- Introduction: R IDE, RStudio, basic data types / objects,
  packages, in-/output, data analysis, visualization

- Business case demo:

  - Extracting 'sales' data from a Nav DB on SQL Server

  - Basic analysis and visualization

  - Advanced visualization using the Shiny framework

- Example: data science going wrong, round-up, resources

- This is an introductory walk-through, no deep dive -
  so no fancy predictions, regression, big data science :-(

History: R - then and now

- Programming language for statistical computing, analysis and visualization, widely used by statisticians, data miners, analysts, data scientists

- Created by **R**oss Ihaka and **R**obert Gentleman, Uni Auckland, in 1993 as an open source implementation of the (1970s) S language

- GNU project, maintained by the R Foundation for Statistical Computing, compiled builds for Mac OS, Linux, Windows, supported by R Consortium

- Extensible through user-created packages, > 13700 available on CRAN

- Commercial support, e.g. since 2007 by Revolution Analytics, acquired by Microsoft in 2015, now provide Microsoft R Open, R Server

- IDEs: R.App, RStudio, R Tools for Visual Studio (deprecated from VS 2019)

- Support for R now in SQL Server, Power BI, Azure ML, Data science VM

Introduction: data objects

- Data types

  - numeric, integer, complex

  - character

  - logical

  - factor

  - Posix types for date/time

  - NA = Not available

- Data structures

  - vector: 1 dim, 1 data type

  - matrix: 2 dim rect, 1 data type

  - list: collection of other objects

  - table: > 2 dimensions

  - data frame
    2 dim rect, cols = vectors

  ▷ `DemoBasics1`

Introduction: packages

- Extensions to the R base system, containing code, data, documentation. Key factor to the success of R; flexible, user contributable. -> CRAN
- `installed.packages()` lists all *installed* packages incl. versions, dependencies, license and other info
- `search()` lists currently *attached* packages
- `install.packages()` downloads and installs packages
- `library()` loads/attaches packages, also `require()`
- Hadley Wickham, chief scientist at RStudio, professor of statistics „Tidyverse": dplyr, tidyr, lubridate, readr, httr, ggplot2 + many more: hadley.nz

▷ DemoBasics2

‹ ›

Introduction: basic data in-/output

- Generic functions `read.table` and `write.table`

  - `read.csv` / `read.csv2` comma/semicolon delimited

  - `read.delim` / `read.delim2` Tab delimited, decimal point/comma

  - `read.fwf` fixed width format

- Some additional I/O packages

  - `reader` functions flexibly load multiple formats <u>fast</u>

  - `foreign` reads data from Minitab, S, SAS, SPSS, Stata, dBase…

  - `DBI/ODBC` database access via ODBC

  - `xlsx` and `readxl` read and write Excel 97/XP/200X files

  - `XML` reads XML and tables from http web sites

Introduction: basic data analysis + visualization

- Analyzing (numeric) data:
  `str()` structure = data types and ranges
  `summary()` Min, max, mean, median, quartiles;
  for factors: count of levels
  `head()`/`tail()` shows top/bottom n rows (default = 6)

- Distribution of values:
  `hist()` shows frequency distribution,
  `boxplot()` for min, max, quartiles, outliers,
 `mosaicplot()` contingency mosaic

  ▷ `DemoBasics3`

Continued… data analysis + visualization

- Libraries: `tidy` for data tidying/reshaping, `ggplot2` implements grammar of graphics, `raster` for geo data

- `apply()` family of functions applies functions to the margins of an array or a matrix

- `gather()/spread()` convert between wide/long format

- `ggplot()` very powerful plot function, plots point, line or bar geometrics etc with versatile parameters

▷ `DemoBasics4`

Business case demo

- We are the distributor for all German petrol stations, with two subsidiaries: NorthTank and SouthFuel
- Business calls „We need some analysis of our 2015 Diesel sales", preferably some visualizations, and „maybe something is wrong…"
- Of interest: distribution by post code zones
- Source: Dynamics Nav ERP database, on the customer card (table „Customer") there's a field called „Sales (LCY)" (= Local currency)
- Publicly available shape- and data files for post code zones

Extracting data & first analysis

- Using ODBC and the DBI package
  (also available: JDBC, RODBC and others)
- `dbConnect()` to establish a connection,
  then `dbGetQuery()` to query the database
- Calculate aggregates (sums) using `ddply()`
- Bar plot: `ggplot() + geom_bar()`
- Line diagram: `ggplot() + geom_line()`

`<>`

Analysis & visualization

- Calculate intervals for sales sums: `cut()`

- libraries `raster`, `rgeos` for visualizing geospatial data

- shapefiles: open vector data format for GIS software, describes points, lines or polygons in these files: .shp shapes, .shx shape index, .dbf attributes, .prj projection

- merge shape and sales data: `merge()`

- plot maps, colouring post code zones according to sales

▷ `DemoTankData`

Use of Shiny framework

- Framework for interactive web applications in R apps consist of `server.R` and `ui.R` or just `app.R`
- `ui` defines screen appearance & controls
- `server` handles any data processing, plotting etc.
- apps can be run in web browser

▷ `DemoShiny/app`

Example: data science going wrong?

- Anscombe's quartet:
- 4 data sets, each with 11 completely different x-y pairs
- yet nearly identical statistical properties
  - Mean of x = 9
  - Mean of y = 7.5
  - Correlation between x and y = 0.816
  - Linear regression y = 3 + 0.5 x

▷ Anscombe

Round-up / conclusions

- With R, a lot is possible in terms of analysis and visualization
- There's probably always a package for that

  But please:
- Know your data
- Look at your data
- Think - does it make sense?
- Consider the influence of outliers
- Don't blindly rely on R 'doing the trick'

Resources online

- https://en.wikipedia.org/wiki/R_(programming_language)

- https://www.r-project.org/ -> Mirrors of CRAN = Comprehensive R Archive Network

- https://www.r-consortium.org/

- http://www.r-bloggers.com/

- www.kdnuggets.com

- www.rseek.org Pimped Google search for R-related subjects

- Twitter hashtag #rstats

- LinkedIn groups R Developers und Users Group, R Programming, The R Project for...

- www.swirlstats.com „Learn R, in R"

- www.coursera.org Data Science specialization (10 courses) MOOC

- www.edx.org

Resources offline

- Beginning R, The statistical programming language
  Dr. Mark Gardener, Wrox/Wiley, ISBN 978-1118164303

- R Cookbook, Paul Teetor, O'Reilly, ISBN 978-0596809157

- R Graphics Cookbook, Winston Chang, O'Reilly,
  ISBN 978-1449316952

- R in a Nutshell, Joseph Adler, O'Reilly, ISBN 978-1449312084

- Practical Data Science with R, Nina Zumel + John Mount,
  Manning publications, ISBN 978-1617291562

Credits

- Titanic data set: www.kaggle.com/c/titanic/data
- SQL Database structure:
  mbs.microsoft.com Dynamics Nav 2016 demo database
- Customer and „sales" data: www.tankerkoenig.de (license CC BY 4.0)
- Shape files:
  - www.suche-postleitzahl.org (Open database license,
  © OpenStreetMap contributors)
  - Bundesamt für Kartographie und Geodäsie, Frankfurt am Main, 2011
- Some icons made by:
  http://www.flaticon.com/authors/hanan (license CC BY 3.0)
- Anscombe's quartet: Francis J. Anscombe 1973

An R primer for SQL folks

Time for some Q & A:

That is: questions that might be of common interest,
and their answers might fit into the remaining time

For slides and scripts: follow link on final slide
or check the SQLBits homepage ;-)

An R primer for SQL folks

Thank you for your interest & keep in touch:

@DerFredo https://twitter.com/DerFredo

de.linkedin.com/in/derfredo

www.xing.com/profile/Thomas_Huetter

Slides and scripts to this presentation will be at
https://github.com/SQLThomas/Conferences/tree/master/Bits2019